

gbtrees

August 10, 2016

```
In [20]: #This code trains the Gradient Boosted Trees for predicting the bike sales
import numpy as np
import pandas as pd
from copy import copy
from sklearn.ensemble import RandomForestRegressor
import csv
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.ensemble import GradientBoostingRegressor

dateparse=lambda x:pd.datetime.strptime(x,'%Y-%m-%d %H:%M:%S')
train=pd.read_csv('train.csv',parse_dates=['datetime'],date_parser=dateparse)
test=pd.read_csv('test.csv',parse_dates=['datetime'],date_parser=dateparse)

#This was required when the number of trees were less than 50. With more trees , almost same p

#test['windspeed']=np.log(test['windspeed']+1)
#train['windspeed']=np.log(train['windspeed']+1)
print train.shape

def extractFeaturesTrain(data):
    #print 'data is ',data
    data['Hour']=data.datetime.dt.hour
    data['DayOfWeek']=data.datetime.dt.dayofweek
    #data['Month']=data.datetime.dt.month
    labels=data['count']
    train_years=data.datetime.dt.year
    train_months=data.datetime.dt.month
    data=data.drop(['datetime','count','casual','registered'], axis = 1)

    return np.array(data),np.array(labels),np.array(train_years),np.array(train_months),(data.

def extractFeaturesTest(data):

    data['Hour']=data.datetime.dt.hour
    data['DayOfWeek']=data.datetime.dt.dayofweek
    #data['Month']=data.datetime.dt.month
    test_years=data.datetime.dt.year
    test_months=data.datetime.dt.month
    data=data.drop(['datetime'], axis = 1)
    return np.array(data),np.array(test_years),np.array(test_months)

train2=copy(train)
```

```

test2=copy(test)
test=np.array(test)
#print 'train2 is ',train2
traind,labelsTrain,train_years,train_months,headers=extractFeaturesTrain(train2)
testd,test_years,test_months=extractFeaturesTest(test2)

submit=np.array((test.shape[0],2))

#train.to_csv('Remodeled Train.csv')
train=np.array(train)
print 'train is \n',traind.shape
print 'labels train are \n',labelsTrain.shape
print 'test is \n',testd.shape

def findLocations(year,month):
    locs=[]
    for i in range(0,test.shape[0]):
        if(test[i][0].year==year and test[i][0].month==month):
            locs.append(i)

    return locs

def findValidDates(year,month):
    locs=[]
    for i in range(0,train.shape[0]):
        if(train[i][0].year<=year and train[i][0].month<=month):
            locs.append(i)

    return locs

'''for i in set(test_years):
    for j in set(test_months):
        print 'Year : ',i,' month ',j:
        testLocs=findLocations(i,j)
        testSubset=testd[testLocs]

        trainLocs=findValidDates(i,j)
        trainSubset=trainind[trainLocs]'''

def findLoss(gold,predicted):
    loss=0
    for i in range(gold.shape[0]):
        loss+=(np.log(predicted[i]+1) -np.log(gold[i]+1))*2

    loss=loss/gold.shape[0]
    #print 'loss is ',loss,' y_pred is ',predicted[i]
    return np.sqrt(loss)

def replaceNegaticeValuesWithZeroAndCountThem(ypred):
    count=0
    for i in range(ypred.shape[0]):
        if(ypred[i]<0):
            ypred[i]=0
            count+=1

```

```

        print 'Number of Negative values predicted are ',count
        return ypred,count

rf=GradientBoostingRegressor(n_estimators=200)
split1=0.8*traind.shape[0]
trainSplit=traind[:split1,:]

testSplit=traind[split1:,:]
labelsSplitTrain=labelsTrain[:split1]
labelsSplitTest=labelsTrain[split1:]
rf.fit(trainSplit,labelsSplitTrain)
ypred=rf.predict(testSplit)
ypred,count=replaceNegaticeValuesWithZeroAndCountThem(ypred)
print 'trainSplit is \n',trainSplit.shape,' and testSplit is \n',testSplit.shape
print 'ypred is \n',ypred
print 'test split is \n',labelsSplitTest
print 'the loss is ',findLoss(labelsSplitTest,ypred)


rf.fit(traind,labelsTrain)
#print 'rf.estimators_ are ',rf.estimators_
print 'testd shape is ',testd.shape
ypred2=rf.predict(testd)
with open('submit2.csv', 'wb') as csvfile:
    resultWriter= csv.writer(csvfile)
    l=['datetime','count']
    resultWriter.writerow(l)
    for i in range(testd.shape[0]):
        #print 'test['',i,'] [0] is ',test[i,0]
        l=[test[i,0],ypred2[i]]
        resultWriter.writerow(l)

allEstimators=rf.estimators_
allEstimators=allEstimators.reshape(1,-1)
allEstimators=allEstimators.tolist()

#print '2 rf.estimators_ are ',allEstimators[0]


importances=rf.feature_importances_
std=np.std([tree.feature_importances_ for tree in allEstimators[0]],axis=0)
indices=np.argsort(importances)[::-1]
print 'Feature Ranking\n'

for f in range(traind.shape[1]):
    print("%d. feature %d %s (%f)" % (f + 1, indices[f],headers[indices[f]], importances[indices[f]]))

fig, ax = plt.subplots()

```

```

ax.set_title('Feature Importances By Gradient Boosted Trees')
ax.bar(range(traind.shape[1]), importances[indices], color="b", yerr=std[indices], align='center')
plt.xticks(range(traind.shape[1]), indices)
ax.set_xlim([-1, traind.shape[1]])
ax.set_xticklabels(headers[indices])
plt.savefig('Feature Importances By Gradient Boosted Trees')
plt.show()

(10886, 12)
train is
(10886, 10)
labels train are
(10886,)
test is
(6493, 10)
Number of Negative values predicted are 32
trainSplit is
(8708, 10) and testSplit is
(2178, 10)
ypred is
[ 34.27567496  35.7852371  158.05560733 ..., 135.21527701 131.96349257
  89.58574943]
test split is
[ 19  19  68 ..., 168 129  88]
the loss is 0.659359778897
testd shape is (6493, 10)
Feature Ranking

1. feature 8 Hour (0.417246)
2. feature 9 DayOfWeek (0.116821)
3. feature 6 humidity (0.096584)
4. feature 2 workingday (0.093433)
5. feature 4 temp (0.085545)
6. feature 5 atemp (0.065820)
7. feature 0 season (0.047824)
8. feature 7 windspeed (0.039049)
9. feature 3 weather (0.024028)
10. feature 1 holiday (0.013651)

/usr/local/lib/python2.7/dist-packages/ipykernel/_main_.py:102: DeprecationWarning: using a non-integer
/usr/local/lib/python2.7/dist-packages/ipykernel/_main_.py:104: DeprecationWarning: using a non-integer
/usr/local/lib/python2.7/dist-packages/ipykernel/_main_.py:105: DeprecationWarning: using a non-integer
/usr/local/lib/python2.7/dist-packages/ipykernel/_main_.py:106: DeprecationWarning: using a non-integer

In [3]: def getTestLocs(year, month):

        locs=[]
        print 'In testlocs year is ', year, ' month is = ', month
        for i in range(0, test.shape[0]):
            if test[i][0].year==year and test[i][0].month==month:
                locs.append(i)
        return locs

In [4]: set(test_years)

```

Out[4]: {2011, 2012}

In [19]: rf2=GradientBoostingRegressor(n_estimators=300)

```
with open('submit3.csv','wb') as csvfile:
    resultWriter=csv.writer(csvfile)
    l=['datetime','count']
    resultWriter.writerow(l)
    for i in set(test_years):
        for j in set(test_months):
            testLocs=getTestLocs(i,j)
            #print 'testLoics are ',testLocs

            testSubset1=testd[testLocs]
            testSubset2=test[testLocs]
            #print 'testSubset2 is ',testSubset2
            trainLocs=np.where(train[:,0]<=min(testSubset2[:,0]))
            trainSubset=train[trainLocs]
            labelsSubset=labelsTrain[trainLocs]
            rf2.fit(trainSubset,labelsSubset)
            ypred3=rf2.predict(testSubset1)
            ypred3,count=replaceNegaticeValuesWithZeroAndCountThem(ypred3)

            for k in range(0,testSubset2.shape[0]):
                l=[testSubset2[k,0],ypred3[k]]
                resultWriter.writerow(l)
```

```
In testlocs year is = 2011 month is = 1
Number of Negative values predicted are 14
In testlocs year is = 2011 month is = 2
Number of Negative values predicted are 4
In testlocs year is = 2011 month is = 3
Number of Negative values predicted are 13
In testlocs year is = 2011 month is = 4
Number of Negative values predicted are 9
In testlocs year is = 2011 month is = 5
Number of Negative values predicted are 13
In testlocs year is = 2011 month is = 6
Number of Negative values predicted are 0
In testlocs year is = 2011 month is = 7
Number of Negative values predicted are 0
In testlocs year is = 2011 month is = 8
Number of Negative values predicted are 0
In testlocs year is = 2011 month is = 9
Number of Negative values predicted are 4
In testlocs year is = 2011 month is = 10
Number of Negative values predicted are 4
In testlocs year is = 2011 month is = 11
Number of Negative values predicted are 6
In testlocs year is = 2011 month is = 12
Number of Negative values predicted are 25
In testlocs year is = 2012 month is = 1
Number of Negative values predicted are 38
In testlocs year is = 2012 month is = 2
```

```

Number of Negative values predicted are 11
In testlocs year is = 2012 month is = 3
Number of Negative values predicted are 16
In testlocs year is = 2012 month is = 4
Number of Negative values predicted are 9
In testlocs year is = 2012 month is = 5
Number of Negative values predicted are 3
In testlocs year is = 2012 month is = 6
Number of Negative values predicted are 8
In testlocs year is = 2012 month is = 7
Number of Negative values predicted are 13
In testlocs year is = 2012 month is = 8
Number of Negative values predicted are 0
In testlocs year is = 2012 month is = 9
Number of Negative values predicted are 8
In testlocs year is = 2012 month is = 10
Number of Negative values predicted are 3
In testlocs year is = 2012 month is = 11
Number of Negative values predicted are 3
In testlocs year is = 2012 month is = 12
Number of Negative values predicted are 36

```

```

In [18]: def getSplits(years,months):
    locsTrain=[]
    locsTest=[]
    for i in range(0,train.shape[0]):
        if (train[i,0].year==years[0] or train[i,0].year==years[1]) and (train[i,0].month
            locsTest.append(i)
        else:
            locsTrain.append(i)

    return locsTrain,locsTest

def getCustomLocsTest(year,month,data):
    locs=[]
    for i in range(0,data.shape[0]):
        if data[i][0].year==year and data[i][0].month==month:
            locs.append(i)
    return locs

def crossValidate():
    months=[12]
    locsTrain,locsTest=getSplits([2011,2012],months)

    testSubset=trainind[locsTest]
    testSubset2=train[locsTest]
    testLabels=labelsTrain[locsTest]
    rf3=GradientBoostingRegressor(n_estimators=300)
    trainSubset=trainind[locsTrain]
    trainSubset2=train[locsTrain]
    trainLabels=labelsTrain[locsTrain]

    for i in [2011,2012]:
        for j in months:

```

```

testLocs=getCustomLocsTest(i,j,testSubset2)
testSubset3=testSubset2[testLocs]
testSubset4=testSubset[testLocs]
testLabels4=testLabels[testLocs]

trainLocs2=np.where(trainSubset2[:,0]<=min(testSubset3[:,0]))

trainSubset3=trainSubset[trainLocs2]
trainLabels3=trainLabels[trainLocs2]
x1=trainSubset2[trainLocs2]
x2=testSubset2[testLocs]

#print 'trainSubset min is ', min(x1[:,0]),' and max is ',max(x1[:,0])
#print 'testSubset min is ', min(x2[:,0]),' and max is ',max(x2[:,0])

#print 'trainSubset3 is ',trainSubset3,'\ntrainLabels3 are ',trainLabels3
rf3.fit(trainSubset3,trainLabels3)
#print 'trained'
ypred=rf3.predict(testSubset4)
ypred,count=replaceNegaticeValuesWithZeroAndCountThem(ypred)
print 'loss with year =',i,' and month = ',j,' is ',findLoss(testLabels4,ypred)

```

```

crossValidate()

```

```

Number of Negative values predicted are 4
loss with year = 2011 and month = 12 is 0.498786860892
Number of Negative values predicted are 7
loss with year = 2012 and month = 12 is 0.552550269169

```

```

In [ ]:

```

```

dataTrain=np.array(train)
dataTrain.shape
plt.plot(dataTrain[:,1],dataTrain[:,11],'*')
plt.show()

```